

DSBDA practical5

May 12, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
[2]: df= pd.read_csv("D:\College Practicals\DSBDApractical5\Social_Network_Ads.csv")
```

```
[3]: df
```

```
[3]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
..
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

[400 rows x 5 columns]

```
[4]: df.isnull()
```

```
[4]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
..
395	False	False	False	False	False
396	False	False	False	False	False
397	False	False	False	False	False
398	False	False	False	False	False
399	False	False	False	False	False

[400 rows x 5 columns]

```
[5]: from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()  
df['Gender'] = le.fit_transform(df['Gender'])
```

```
[6]: df
```

```
[6]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	1	19	19000	0
1	15810944	1	35	20000	0
2	15668575	0	26	43000	0
3	15603246	0	27	57000	0
4	15804002	1	19	76000	0
..
395	15691863	0	46	41000	1
396	15706071	1	51	23000	1
397	15654296	0	50	20000	1
398	15755018	1	36	33000	0
399	15594041	0	49	36000	1

[400 rows x 5 columns]

```
[7]: # Select the columns you want to include in the covariance matrix  
columns = ['Age', 'EstimatedSalary', 'Purchased']
```

```
# Create a new DataFrame that contains only the selected columns  
df_selected = df[columns]
```

```
# Build the covariance matrix  
covariance_matrix = df_selected.cov()
```

```
# Print the covariance matrix  
print(covariance_matrix)
```

	Age	EstimatedSalary	Purchased
Age	109.890702	5.548738e+04	3.131165
EstimatedSalary	55487.380952	1.162603e+09	5924.367168
Purchased	3.131165	5.924367e+03	0.230269

```
[8]: # Select the independent variables  
X = df.drop('Purchased', axis=1)
```

```
# Select the dependent variable  
Y = df['Purchased']
```

```
[9]: from sklearn.model_selection import train_test_split
```

```
[10]: # Select the independent and dependent variables
X = df.drop('Purchased', axis=1)
Y = df['Purchased']
```

```
[22]: # Split the dataset into a training set and a testing set
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3,
↳random_state=42)
```

```
[23]: #import the class
from sklearn.linear_model import LogisticRegression
```

```
[24]: #instantiate the model

logreg = LogisticRegression()
```

```
[25]: logreg.fit(X_train,Y_train)
```

```
[25]: LogisticRegression()
```

```
[30]: y_pred=logreg.predict(X_test)
```

```
[32]: print("Predicted values for the testing data: ", Y_pred)
```

```
Predicted values for the testing data: [0 0 0 0 1 0 0 0 0 0 0 0 0 1 0 1 0 1 0 0
0 1 0 0 0 0 1 0 0 0 1 1 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 1 1 0 1 0 1 0 0 0 0 0 0
0 0 0 0 0 1 1 0 0 0 0 1 0 0 0 1 1 0 1 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 0
0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 0 1 0 0 0 0 0
0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 0 1 0 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 1 1 0 0
1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 0 0 0 1 0 1 1 0 0 0 0 0
0 1 0 0 0 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 1 1 0 0 0 1 0 1 0 1 0 0 0 0]
```

```
[27]: # Predict the Y values for the training data
Y_pred = logreg.predict(X_train)
```

```
[29]: print("Predicted values for the training data: ", Y_pred)
```

```
Predicted values for the training data: [0 0 0 0 1 0 0 0 0 0 0 0 0 1 0 1 0 1 0
0 0 1 0 0 0 0 1 0 0 0 1 1 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 1 1 0 1 0 1 0 0 0 0 0
0 0 0 0 0 1 1 0 0 0 0 1 0 0 0 1 1 0 1 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 0
0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 0 1 0 0 0 0 0
0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 0 1 0 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 1 1 0 0
1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 0 0 0 1 0 1 1 0 0 0 0 0
0 1 0 0 0 1 0 0 0 0 1 0 1 0 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 1 1 0 0 0 1 0 1 0 1 0 0 0 0]
```

```
[35]: from sklearn.metrics import accuracy_score

      # Generate predicted values for the training data
      Y_pred = logreg.predict(X_train)
```

```
[36]: # Calculate the accuracy of the model
      train_accuracy = accuracy_score(Y_train, Y_pred)
```

```
[37]: # Print the accuracy
      print("Accuracy on training data: {:.2f}%".format(train_accuracy * 100))
```

Accuracy on training data: 77.50%

```
[40]: from sklearn.metrics import precision_score, confusion_matrix, recall_score

      # Calculate the precision score of the model
      train_precision = precision_score(Y_train, Y_pred)
```

```
[41]: print("Precision on training data: {:.2f}%".format(train_precision * 100))
```

Precision on training data: 84.75%

```
[43]: # Calculate the confusion matrix of the model
      train_cm = confusion_matrix(Y_train, Y_pred)
```

```
[44]: print("Confusion matrix on training data:\n", train_cm)
```

Confusion matrix on training data:
[[167 9]
 [54 50]]

```
[46]: # Calculate the recall score of the model
      train_recall = recall_score(Y_train, Y_pred)
```

```
[47]: print("Recall on training data: {:.2f}%".format(train_recall * 100))
```

Recall on training data: 48.08%

```
[ ]:
```